







Extra Challenges and Questions Solutions

- When you use ksqlDB to join streaming data, you must ensure that **your streams and tables are co-partitioned**. To be considered co-partitioned we look at 3 properties:
 - Records have the same keying scheme - both topics are keyed on `driver-id`.
 - Records have the same number of partitions - both topics have 3 partitions.
 - Records have the same partitioning strategy - the producers and Kafka Connect are using the default murmur2 hash partitioner.

These properties guarantee that records of the same key will always be delivered to the same partition number in both topics. For example, `driver-3` will be delivered to partition 1 for both `driver-profiles-ksql` and `driver-positions-avro`.

